


REGRESSION ANALYSIS AND RESAMPLING METHODS

FYS-STK4155: PROJECT 1

Morten Ledum & Håkon Kristiansen

 github.com/mortele/FYS-STK4155

September 27, 2018

Abstract

We parameterize digital terrain data using linear regression analysis algorithms: Ordinary least squares (OLS), Ridge regression, and Lasso regression. The bootstrap resampling technique is used to gauge the bias and variance of the models. We use basis sets of homogeneous monomials in two variables, up to and including total degree 5. We find that xxxxxx.

For initial validation of our models, we employ the test function of R. Franke.^{franke1979critical}

Contents

I	Introduction	2
II	Theory	2
A	Linear regression	2
B	Ordinary least squares	2
	The design matrix	3
C	Polynomial basis sets	3
D	Ridge regression	3
	Singular-value decomposition	4
	Ridge regression using the SVD	4
E	Lasso regression	5
F	Principal components	5
G	Assessing model accuracy	6
H	Resampling and the <i>Bootstrap</i> method	6
III	Data sets	7
I	The Franke function	7
J	Terrain data	7
IV	Results and discussion	10
K	Verification of the models: The Franke function	10
L	Terrain data parametrization	10
V	Conclusion	10

INTRODUCTION

Nulla malesuada porttitor diam. Donec felis erat, congue non, volutpat at, tincidunt tristique, libero. Vivamus viverra fermentum felis. Donec nonummy pellentesque ante. Phasellus adipiscing semper elit. Proin fermentum massa ac quam. Sed diam turpis, molestie vitae, placerat a, molestie nec, leo. Maecenas lacinia. Nam ipsum ligula, eleifend at, accumsan nec, suscipit a, ipsum. Morbi blandit ligula feugiat magna. Nunc eleifend consequat lorem. Sed lacinia nulla vitae enim. Pellentesque tincidunt purus vel magna. Integer non enim. Praesent euismod nunc eu purus. Donec bibendum quam in tellus. Nullam cursus pulvinar lectus. Donec et mi. Nam vulputate metus eu enim. Vestibulum pellentesque felis eu massa.

THEORY

In the following we briefly introduce the theory underlying the technical aspects of the present work. We begin by considering linear regression in general, and the ordinary least squares (OLS) method.

Linear regression

In order to introduce the least squares methods, we consider a case in which p characteristics of n samples are measured. The outcome, or the *response*, is denoted \mathbf{y} : a vector of size n . The measured characteristics, denoted the predictors, are organized in a matrix \mathbf{X} of size $n \times p$. This is called the *design matrix*.

In regression analysis, we aim to explain the response in terms of the predictors, i.e. construct a function $\mathbf{y}(\mathbf{X})$. Assuming a linear relationship between \mathbf{X} and \mathbf{y} gives rise to *linear regression*, in which the response can be written as

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\varepsilon}, \quad (1)$$

where $\boldsymbol{\varepsilon}$ denotes the deviation of the linear model $\mathbf{X}\boldsymbol{\beta}$ and the response \mathbf{y} and $\boldsymbol{\beta}$ is a parameter vector containing the linear regression coefficients β_i . The β_i variables are the unknowns in the linear regression problem, and they represent the partial derivative of the *modelled* response w.r.t. the descriptors.

In any non-trivial case, the error terms ε_i in the error vector $\boldsymbol{\varepsilon}$ will be non-zero. In this case, we regard our linear ansatz as a *model* of

the true response, the observed values y_i . We denote our model by $\tilde{\mathbf{y}}$, and define

$$\begin{aligned} \tilde{\mathbf{y}} &= \mathbf{X}\boldsymbol{\beta} \\ &= \mathbf{y} - \boldsymbol{\varepsilon}. \end{aligned} \quad (2)$$

The objective of linear regression thus emerges: Determine $\boldsymbol{\beta}$ in such a way that $\boldsymbol{\varepsilon}$ is minimized, thus giving a best possible linear fit of the response (minimizing the deviation $|\mathbf{y} - \tilde{\mathbf{y}}|$).

Ordinary least squares

In order to *minimize* the error $\boldsymbol{\varepsilon}$, we must define exactly what that means. We require a functional expression—commonly referred to as the *cost function*—and a metric in which to calculate it's size. Choosing the Euclidean L^2 norm ($\|\mathbf{v}\|_2 = \sqrt{\sum_i v_i^2}$) and the absolute value of $\boldsymbol{\varepsilon}$ as the metric and cost function, respectively, leads to the *ordinary least squares* (OLS) method. Defining the cost function,

$$\begin{aligned} C(\boldsymbol{\beta}) &= \|\mathbf{y} - \tilde{\mathbf{y}}\|_2^2 \\ &= \|\mathbf{y} - \mathbf{X}\boldsymbol{\beta}\|_2^2 \\ &= \sum_{i=1}^n \left| y_i - \beta_0 - \sum_{j=1}^p X_{ij}\beta_j \right|^2, \end{aligned} \quad (3)$$

we can formulate the OLS method as computing $\boldsymbol{\beta}_{\text{optimal}}$ by

$$\boldsymbol{\beta}_{\text{optimal}} = \arg \min_{\boldsymbol{\beta}} \{C(\boldsymbol{\beta})\}. \quad (4)$$

In order to find $\boldsymbol{\beta}_{\text{optimal}}$, we may simply differentiate $C(\boldsymbol{\beta})$ w.r.t. $\boldsymbol{\beta}$ and enforce $\partial C(\boldsymbol{\beta})/\partial \boldsymbol{\beta} = 0$. Following Hastie, Tibshirani & Friedman, ^{trevor2009elements} we find that

$$\begin{aligned} \frac{\partial C(\boldsymbol{\beta})}{\partial \boldsymbol{\beta}} &= \frac{\partial}{\partial \boldsymbol{\beta}} (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})^T (\mathbf{y} - \mathbf{X}\boldsymbol{\beta}) \\ &= -2\mathbf{X}^T (\mathbf{y} - \mathbf{X}\boldsymbol{\beta}) \stackrel{!}{=} 0 \\ \Rightarrow \mathbf{X}^T \mathbf{y} &= \mathbf{X}^T \mathbf{X} \boldsymbol{\beta} \\ \boldsymbol{\beta}_{\text{optimal}} &= (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}, \end{aligned} \quad (5)$$

where we have written $\|\mathbf{y} - \mathbf{X}\boldsymbol{\beta}\|_2^2$ as $(\mathbf{y} - \mathbf{X}\boldsymbol{\beta})^T (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})$. We note that even though $\mathbf{X} \in \mathbb{R}^{n \times p}$ is a “large” matrix (assuming the number of observations $n \gg p$ the number of predictors per observation), the product $\mathbf{X}^T \mathbf{X} \in \mathbb{R}^{p \times p}$ is “small”. Thus explicitly inverting $\mathbf{X}^T \mathbf{X}$ is not a problem on a modern computer.

We note that the *model prediction* may now be calculated simply as $\tilde{\mathbf{y}} = \mathbf{X}\boldsymbol{\beta}_{\text{optimal}}$.

This represents the optimal linear model subject under the Euclidean norm of the cost function as given in Eq. (3). This method was first rigorously described by Legendre in 1805.^{legendre1805nouvelles}

The design matrix

The design matrix, \mathbf{X} , can in principle contain any set of linearly independent functions of the predictors¹. Every column in the design matrix corresponds to a mapping of the predictors, with elements $\mathbf{p}_i \mapsto \mathbf{X}_{ij}$. We will now consider an example of such a design matrix. We use two predictors—we will denote them x and y —with the response y . We introduce our model using the mappings $(x, y) \mapsto x$, $(x, y) \mapsto y$, and $(x, y) \mapsto xy$. Including also what is commonly referred to as the intercept, this gives rise to the design matrix

$$\mathbf{X} = \begin{bmatrix} 1 & x_1 & y_1 & x_1 y_1 \\ 1 & x_2 & y_2 & x_2 y_2 \\ & & \vdots & \\ 1 & x_{n-1} & y_{n-1} & x_{n-1} y_{n-1} \\ 1 & x_n & y_n & x_n y_n \end{bmatrix}. \quad (6)$$

For inputs (x_i, y_i) our model $\tilde{\mathbf{y}}$ now returns

$$\begin{aligned} \tilde{\mathbf{y}}_i &= \mathbf{x}_i^T \boldsymbol{\beta} \\ &= \beta_0 + \beta_1 x_i + \beta_2 y_i + \beta_3 x_i y_i, \end{aligned} \quad (7)$$

where we used the shorthand notation \mathbf{x}_i^T to denote $\text{Row}_i(\mathbf{X})$.

Before we continue describing the Ridge and Lasso regression schemes, we briefly introduce the basis sets used in this project.

Polynomial basis sets

Throughout the present work we employ a basis set of homogeneous monomials². We will be working with 2D terrain data, and thus will need to consider monomials of up to and including two variables— x and y —in all possible homogeneous combinations. Disregarding the zero degree monomial, there are two possible such terms of degree up to and including one. These are simply x and y . Moving up to degree two, we must include x^2 , y^2 , and xy , for a total of

five terms up to and including degree 2. Degree three adds an additional four terms: x^3 , $x^2 y$, $x y^2$, and y^3 , and so on. In general, there are $n + 1$ such terms for monomials of degree n , namely

$$x^n, x^{n-1}y, x^{n-2}y^2, \dots, xy^{n-1}, \text{ and } y^n.$$

The total basis sets of all such monomials of degree *up to and including* degree n — \mathcal{B}_n —thus contains

$$\text{size}(\mathcal{B}_n) \sum_{k=2}^{n+1} k = \frac{n(n+3)}{2} \quad (8)$$

terms.

Ridge regression

As mentioned in section B, defining exactly what we mean by *minimizing the error* requires a cost function and a metric. The previous choice of $C(\boldsymbol{\beta}) = \|\mathbf{y} - \tilde{\mathbf{y}}\|_2^2$ is obviously not the only possible one. In fact, a possibly superior method may be devised by keeping the Euclidean L^2 norm, but including a term in the cost function which penalizes large values of β_i . Such an approach was first used in statistics by Hoerl & Kennard,^{hoerl1970ridge} but was proposed already in the 1940s by Andrey Tikhonov.^{tikhonov1943stability} Taking the cost function to be

$$C_T(\boldsymbol{\beta}) = \|\mathbf{y} - \tilde{\mathbf{y}}\|_2^2 + \|\boldsymbol{\Gamma} \boldsymbol{\beta}\|_2^2 \quad (9)$$

gives rise to a regressions scheme known as Tikhonov regularization. The Tikhonov matrix $\boldsymbol{\Gamma}$ governs the form of the regularization term. A simple case of $\boldsymbol{\Gamma} = \sqrt{\lambda} \mathbf{1}$, favoring solutions with small (in the L^2 norm sense) values of the parameters $\boldsymbol{\beta}$, results in the *Ridge regression* of Hoerl & Kennard. The $\lambda \geq 0$ parameter here represents a tuneable penalty for large $\boldsymbol{\beta}$ values. We note that $\lambda = 0$ recovers the OLS method of section B.

Writing out the cost function of Ridge re-

¹We require linear independence to ensure the normal equations have a unique solution

²A homogeneous polynomial is a polynomial in which all terms have the same total degree, e.g. $xy + y^2 + x^2$ is a homogeneous polynomial, x is a homogeneous monomial, while $xy + x^3$ is *not*. Monomials are simply polynomials with only a single term.

gression, we find that

$$\begin{aligned}
C_R(\beta) &= \|\mathbf{y} - \tilde{\mathbf{y}}\|_2^2 + \lambda \|\beta\|_2^2 \\
&= \|\mathbf{y} - \mathbf{X}\beta\|_2^2 + \lambda \|\beta\|_2^2 \\
&= \sum_{i=1}^n \left| y_i - \beta_0 - \sum_{j=1}^p X_{ij}\beta_j \right|^2 + \lambda \sum_{j=1}^p \beta_j^2,
\end{aligned} \tag{10}$$

where we have left out the intercept β_0 from the regularization term. This is done to ensure the solutions do not explicitly depend on the zero point chosen for \mathbf{y} , i.e. adding a constant to each response value y_i would not result in a simple shift of the predictions by the same amount. ^{trevor2009elements}

In the same way as before, we may differentiate the cost function and enforce $C_R(\beta) = 0$ in order to find the optimal β . We obtain

$$\begin{aligned}
\frac{\partial C_R(\beta)}{\partial \beta} &= \frac{\partial}{\partial \beta} \left[(\mathbf{y} - \mathbf{X}\beta)^T (\mathbf{y} - \mathbf{X}\beta) + \lambda \mathbf{1}^T \beta \right] \\
&= -2\mathbf{X}^T (\mathbf{y} - \mathbf{X}\beta) + 2\lambda \mathbf{1} \beta \stackrel{!}{=} 0 \\
\Rightarrow \mathbf{X}^T \mathbf{y} &= \mathbf{X}^T \mathbf{X} \beta + \lambda \mathbf{1} \beta \\
\beta_{\text{optimal}}^R &= (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{1})^{-1} \mathbf{X}^T \mathbf{y}.
\end{aligned} \tag{11}$$

Singular-value decomposition

Before we continue, we introduce briefly the singular-value decomposition (SVD). Note carefully that *any* $m \times n$ matrix \mathbf{A} can be decomposed into a product like this.

Let \mathbf{A} be an $m \times n$ matrix with rank r . Then there exists an $m \times n$ matrix Σ in which the first r diagonal entries are the singular values of \mathbf{A} , $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_r > 0$, and there exist an $m \times m$ orthogonal matrix \mathbf{U} and an $n \times n$ orthogonal matrix \mathbf{V} such that

$$\mathbf{A} = \mathbf{U} \Sigma \mathbf{V}^T. \tag{12}$$

All entries in Σ outside of the first r diagonal elements are zero. The singular values of \mathbf{A} denote the square roots of the eigenvalues of $\mathbf{A}^T \mathbf{A}$. ^{lay2012linear}

The (first r columns of the) matrix \mathbf{U} contains the eigenvectors of $\mathbf{A}^T \mathbf{A}$ and represents an orthonormal basis for the column space of \mathbf{A} , $\text{Col } \mathbf{A}$. The (first r columns of the) matrix \mathbf{V} contains the eigenvectors of $\mathbf{A} \mathbf{A}^T$ and represents an orthonormal basis for the row space

of \mathbf{A} , $\text{Row } \mathbf{A}$. The remaining $m - r$ and $n - r$ columns of \mathbf{U} and \mathbf{V} form orthonormal bases for $\text{Nul } \mathbf{A}$ and $\text{Nul } \mathbf{A}^T$.

Thus we can interpret the SVD loosely as finding the orthonormal bases of $\text{Col } \mathbf{A}$ and $\text{Row } \mathbf{A}$ such that application of \mathbf{A} maps $\mathbf{v}_i \mapsto \sigma_i \mathbf{u}_i$.

Ridge regression using the SVD

If we consider the SVD of \mathbf{X} , $\mathbf{X} = \mathbf{U} \Sigma \mathbf{V}^T$, and compute $\mathbf{X}^T \mathbf{X}$ we find that

$$\begin{aligned}
\mathbf{X}^T \mathbf{X} &= (\mathbf{U} \Sigma \mathbf{V}^T)^T (\mathbf{U} \Sigma \mathbf{V}^T) \\
&= \mathbf{V} \Sigma^T \mathbf{U}^T \mathbf{U} \Sigma \mathbf{V}^T \\
&= \mathbf{V} \Sigma^T \Sigma \mathbf{V}^T = \mathbf{V} \Sigma^2 \mathbf{V}^T,
\end{aligned} \tag{13}$$

where the orthogonality of \mathbf{U} made $\mathbf{U}^T \mathbf{U} = \mathbf{1}$. Inserting this into the expression for the optimal β_{optimal}^R (Eq. (11)) yields ^{trevor2009elements}

$$\begin{aligned}
\beta_{\text{optimal}}^R &= (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{1})^{-1} \mathbf{X}^T \mathbf{y} \\
&= (\mathbf{V} \Sigma^2 \mathbf{V}^T + \lambda \mathbf{1})^{-1} \mathbf{X}^T \mathbf{y} \\
&= (\mathbf{V} \Sigma^2 \mathbf{V}^T + \lambda \mathbf{1} \mathbf{V} \mathbf{V}^T)^{-1} \mathbf{X}^T \mathbf{y} \\
&= (\mathbf{V} \Sigma^2 \mathbf{V}^T + \lambda \mathbf{V} \mathbf{1} \mathbf{V}^T)^{-1} \mathbf{X}^T \mathbf{y},
\end{aligned}$$

where we multiplied by $\mathbf{1} = \mathbf{V} \mathbf{V}^T$ (recall that $\mathbf{V}^{-1} = \mathbf{V}^T$ due to orthogonality) and used the fact that the identity matrix commutes with any other matrix. Furthermore, we find

$$\begin{aligned}
\beta_{\text{optimal}}^R &= (\mathbf{V} [\Sigma^2 + \lambda \mathbf{1}] \mathbf{V}^T)^{-1} \mathbf{X}^T \mathbf{y} \\
&= \mathbf{V} [\Sigma^2 + \lambda \mathbf{1}]^{-1} \mathbf{V}^T \mathbf{X}^T \mathbf{y} \tag{14} \\
&= \mathbf{V} [\Sigma^2 + \lambda \mathbf{1}]^{-1} \mathbf{V}^T (\mathbf{U} \Sigma \mathbf{V}^T)^T \mathbf{y} \\
&= \mathbf{V} [\Sigma^2 + \lambda \mathbf{1}]^{-1} \mathbf{V}^T \mathbf{V} \Sigma^T \mathbf{U}^T \mathbf{y} \\
&= \mathbf{V} [\Sigma^2 + \lambda \mathbf{1}]^{-1} \Sigma \mathbf{U}^T \mathbf{y},
\end{aligned} \tag{15}$$

where we used that Σ is diagonal, so $\Sigma^T = \Sigma$. Note that since $(\mathbf{A} \mathbf{B})^{-1} = \mathbf{B}^{-1} \mathbf{A}^{-1}$, we can rewrite the inverse product

$$\begin{aligned}
(\mathbf{V} [\Sigma^2 + \lambda \mathbf{1}] \mathbf{V}^T)^{-1} &= (\mathbf{V}^T)^{-1} [\Sigma^2 + \lambda \mathbf{1}]^{-1} \mathbf{V}^{-1} \\
&= \mathbf{V} [\Sigma^2 + \lambda \mathbf{1}]^{-1} \mathbf{V}^T,
\end{aligned}$$

as was done in Eq. (14). Since we are now taking the inverse of a diagonal matrix, we can simply write out the terms. Note that the inverse will itself be diagonal, and given by

$$([\Sigma^2 + \lambda \mathbf{1}]^{-1})_{ii} = \frac{1}{\sigma_{ii}^2 + \lambda}. \tag{16}$$

Rewriting the OLS scheme in terms of the SVD yields a very similar

$$\boldsymbol{\beta}_{\text{optimal}} = \mathbf{V} (\boldsymbol{\Sigma}^2)^{-1} \mathbf{V}^T \boldsymbol{\Sigma}^T \mathbf{U}^T \mathbf{y}. \quad (17)$$

Lasso regression

As mentioned repeatedly, we are free to choose what we mean by *minimizing the error* w.r.t. what metric and what cost function we use. Whereas the Ridge regression of section D used L^2 regularization by adding a $\lambda \|\boldsymbol{\beta}\|_2^2$ term to $C(\boldsymbol{\beta})$, we may instead try a L^1 regularization. This constitutes setting up the cost function as

$$\begin{aligned} C_L(\boldsymbol{\beta}) &= \|\mathbf{y} - \tilde{\mathbf{y}}\|_2^2 + \lambda \|\boldsymbol{\beta}\|_1 \\ &= \|\mathbf{y} - \mathbf{X}\boldsymbol{\beta}\|_2^2 + \lambda \|\boldsymbol{\beta}\|_1 \\ &= \sum_{i=1}^n \left| y_i - \beta_0 - \sum_{j=1}^p X_{ip} \beta_p \right|^2 + \lambda \sum_{j=1}^p |\beta_j|^2. \end{aligned} \quad (18)$$

Originally popularized by Tibshirani, the *Lasso regression* has certain potential advantages over the OLS and Ridge regression schemes. Most notably, the Lasso can perform *variable selection*, i.e. some β_j s may be identically zero as a result of the minimization. The name Lasso is short for “least absolute shrinkage and selection operator”.

Computing the derivative of the Lasso cost function yields

$$\begin{aligned} \frac{\partial C_L(\boldsymbol{\beta})}{\partial \beta_j} &= \frac{\partial C_{\text{OLS}}(\boldsymbol{\beta})}{\partial \beta_j} + \lambda \sum_{j=1}^p \frac{\partial}{\partial \beta_j} |\beta_j| \\ &= \frac{\partial C_{\text{OLS}}(\boldsymbol{\beta})}{\partial \beta_j} + \lambda \frac{\beta_j}{\sqrt{\beta_j^2}} \\ &= \frac{\partial C_{\text{OLS}}(\boldsymbol{\beta})}{\partial \beta_j} + \lambda \text{sgn}(\beta_j) \stackrel{!}{=} 0. \end{aligned} \quad (19)$$

In general, this can not be directly solved for $\boldsymbol{\beta}_{\text{optimal}}^L$ as in the case of OLS or the Ridge scheme. Under the assumption that \mathbf{X} is orthogonal, an explicit solution exists and is given by

$$(\boldsymbol{\beta}_{\text{optimal}}^L(\lambda))_j = \text{sgn}(\beta_j^{\text{OLS}}) (|\beta_j^{\text{OLS}}| - \lambda)_+, \quad (20)$$

where $(\cdot)_+$ represents the positive part of \cdot . In the general case, an iterative solver must be used to compute $\boldsymbol{\beta}_{\text{optimal}}^L$.

Principal components

The following section follows Hastie, Tibshirani & Friedman.

Recall from section that the SVD matrix \mathbf{U} represents an orthonormal basis for the column space of the decomposed matrix \mathbf{A} . If we consider the prediction resulting from OLS, and perform a SVD of \mathbf{X} we find that

$$\begin{aligned} \tilde{\mathbf{y}}_{\text{OLS}} &= \mathbf{X}\boldsymbol{\beta}_{\text{OLS}} \\ &= \mathbf{X} (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} \\ &= \mathbf{X} [(\mathbf{U}\boldsymbol{\Sigma}\mathbf{V}^T)^T \mathbf{U}\boldsymbol{\Sigma}\mathbf{V}^T]^{-1} \mathbf{X}^T \mathbf{y} \\ &= \mathbf{U}\boldsymbol{\Sigma}\mathbf{V}^T [\mathbf{V}\boldsymbol{\Sigma}^2\mathbf{V}^T]^{-1} \mathbf{V}\boldsymbol{\Sigma}^T \mathbf{U}^T \mathbf{y} \\ &= \mathbf{U}\boldsymbol{\Sigma}\mathbf{V}^T \mathbf{V} [\boldsymbol{\Sigma}^2]^{-1} \mathbf{V}^T \mathbf{V}\boldsymbol{\Sigma}^T \mathbf{U}^T \mathbf{y} \\ &= \mathbf{U}\boldsymbol{\Sigma} [\boldsymbol{\Sigma}^2]^{-1} \mathbf{V}^T \mathbf{V}\boldsymbol{\Sigma}^T \mathbf{U}^T \mathbf{y} \\ &= \mathbf{U}\boldsymbol{\Sigma} [\boldsymbol{\Sigma}^2]^{-1} \boldsymbol{\Sigma}^T \mathbf{U}^T \mathbf{y} \\ &= \mathbf{U}\mathbf{U}^T \mathbf{y}. \end{aligned} \quad (21)$$

Assuming the derivation for $\tilde{\mathbf{y}}_R$ for the Ridge regression yields

$$\begin{aligned} \tilde{\mathbf{y}}_R &= \mathbf{X}\boldsymbol{\beta}_R \\ &= \mathbf{X} (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y} \\ &= \mathbf{U}\boldsymbol{\Sigma} (\boldsymbol{\Sigma}^2 + \lambda \mathbf{I})^{-1} \boldsymbol{\Sigma}^T \mathbf{U}^T \mathbf{y} \\ &= \sum_{j=1}^p \mathbf{u}_j \frac{\sigma_j^2}{\sigma_j^2 + \lambda} \mathbf{u}_j^T \mathbf{y}, \end{aligned} \quad (22)$$

where σ_j denotes the diagonal elements of the diagonal matrix $\boldsymbol{\Sigma}$, i.e. $(\boldsymbol{\Sigma})_{jj} = \sigma_j$. Note now that $\mathbf{U}^T \mathbf{y}$ are the coordinates of \mathbf{y} w.r.t. the orthonormal basis \mathbf{U} . Comparing Eqs. (21) and (22), we note that the coordinates of \mathbf{y} in both cases are computed in the orthonormal basis of $\text{Col } \mathbf{X}$ (as specified by the SVD matrix \mathbf{U}), but the Ridge scheme also *shrinks* the coordinates. The shrinkage is large whenever σ_j^2 is small. Recalling that $\mathbf{X}^T \mathbf{X} = \mathbf{V}\boldsymbol{\Sigma}^2\mathbf{V}^T$ is an eigendecomposition of $\mathbf{X}^T \mathbf{X}$, with \mathbf{V} containing the eigenvectors, we denote these eigenvectors \mathbf{v}_j to be the *principal components* of \mathbf{X} . The eigenvalues contained in $\boldsymbol{\Sigma}^2$ are precisely the proportionality factors, σ_j^2 , involved in the shrinkage.

The diagonalization of $\mathbf{X}^T \mathbf{X}$ constitutes a coordinate transform into a coordinate system in which $\mathbf{X}^T \mathbf{X}$ itself is diagonal. The matrix $\mathbf{X}^T \mathbf{X}$ is the covariance matrix (apart from a constant factor $1/N$). In the orthonormal basis of \mathbf{V} , the covariance matrix is diagonal and

contains the variances σ_j^2 of these linear combinations of the columns of \mathbf{X} .

The first principal component direction has the property that $\mathbf{z}_1 = \mathbf{X}\mathbf{v}_1$ has the largest variance of all the linear combinations of the columns of \mathbf{X} . In general, the principal components are ordered such that $\text{Var } \mathbf{z}_1 \geq \text{Var } \mathbf{z}_2 \geq \dots \geq \text{Var } \mathbf{z}_N$.

In essence, the first principal component represent the direction in $\text{Col } \mathbf{X}$ in which the variance is highest, the second principal component represents the corresponding direction in which the variance is highest, apart from the first, and so on. It is clear that the Ridge regression scheme simply rotates the OLS solution into the principal components of the design matrix, and then shrinks the coefficients corresponding to low-variance components.

Assessing model accuracy

The goal in machine learning is to use a chosen model to make predictions. In particular, we need to quantify how close a predicted response value for a given observation is to the true response value. In the regression setting, a common measure is the *mean squared error* (MSE), given by

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n \left(y_i - \hat{f}(x_i) \right)^2, \quad (23)$$

where $\hat{f}(x_i)$ is the prediction for the i th observation y_i . Computing the MSE on the training set result in what we could call the training MSE and gives a measure of how well the model fits the training data. However, we are mainly interested in how good predictions are on previously unseen data.

Suppose that we have set of test data, $\{x_i, y_i\}$, which is not used during training. Then we could compute the MSE on the test set, giving a measure on how well the method works on unseen data. We would like to select the method which has the lowest test MSE.

The challenge now is how do we estimate the test MSE? Of course, if test data are available we just compute. The problem is that test data are not available in general. So if do not have test data at hand, a natural approach seems to choose the method that minimizes the training MSE. However, there is no guarantee that low training MSE results in the lowest test MSE. As an extreme example suppose you draw a curve

that passes through every point in a data set, this would result in perfect train MSE. If you try to use this model on another data set where I move one datapoint significantly it would result in a high test MSE.

One can show that the expected test MSE, for a given value x_0 , can always be decomposed as

$$E[(y_0 - \hat{f}(x_0))]^2 = \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{Var}(\epsilon). \quad (24)$$

Here $E[(y_0 - \hat{f}(x_0))]^2$ defines the expected test MSE, and refers to the average test MSE that we would obtain if we repeatedly estimated f using a large number of training sets, and tested each at x_0 .

The variance, $\text{Var}(\hat{f}(x_0))$, measures how much \hat{f} would change if we trained the model on a different data set. Ideally \hat{f} should not vary too much between training sets. Generally more flexible methods have higher variance, which can give large deviations in \hat{f} when changes in the training set is small.

Bias measures the the error introduced by assuming for example a linear relationship between the response \mathbf{y} and the predictors \mathbf{X} . Thus, if the true relation is non-linear this will result in bias. In particular, if the data is highly non-linear it would be impossible to produce accurate estimates by assuming a linear relationship between the response and predictors. Generally more flexible methods will have less bias.

Equation 24 suggests that we want both low variance and low bias in order to minimize the expected test error. The above discussion implies that these are conflicting interests.

Resampling and the *Bootstrap* method

Nulla malesuada porttitor diam. Donec felis erat, congue non, volutpat at, tincidunt tristique, libero. Vivamus viverra fermentum felis. Donec nonummy pellentesque ante. Phasellus adipiscing semper elit. Proin fermentum massa ac quam. Sed diam turpis, molestie vitae, placerat a, molestie nec, leo. Maecenas lacinia. Nam ipsum ligula, eleifend at, accumsan nec, suscipit a, ipsum. Morbi blandit ligula feugiat magna. Nunc eleifend consequat lorem. Sed lacinia nulla vitae enim. Pellentesque tincidunt purus vel magna. Integer non enim. Praesent euismod nunc eu purus. Donec bibendum quam in tellus. Nullam cursus pulvinar lectus. Donec et mi. Nam vulputate metus eu enim. Vestibulum

pellentesque felis eu massa.

DATA SETS

We are chiefly interested in parametrizing digital terrain data. However, in order to test and validate our implementation of the regression model and the resampling technique, we employ the Franke function ^{franke1979critical} as a test case before considering real data.

The Franke function

The test function of Franke—originally developed to test and rate different surface interpolation techniques—is “a surface with a variety of behaviour” which consists of “two Gaussian peaks and a sharper Gaussian dip superimposed on a surface sloping towards the first quadrant.” ^{franke1979critical} It is noted by Franke in the his original paper that the slope was introduced mainly as a visual aid and presumably had little impact on the actual interpolations performed.

More specifically, the Franke function $f_F(x, y)$ takes the full form

$$\begin{aligned} f_F(x, y) = & \frac{3}{4} \exp \left\{ \frac{-1}{4} \left[(9x - 2)^2 + (9y - 2)^2 \right] \right\} \\ & + \frac{3}{4} \exp \left\{ \frac{-1}{49} (9x + 1)^2 + \frac{1}{10} (9y + 1)^2 \right\} \\ & + \frac{1}{2} \exp \left\{ \frac{-1}{4} \left[(9x - 7)^2 + (9y - 3)^2 \right] \right\} \\ & - \frac{1}{5} \exp \left\{ \frac{-1}{4} \left[(9x + 4)^2 + (9y - 7)^2 \right] \right\}. \end{aligned} \quad (25)$$

A plot of the $f_F(x, y)$ surface can be seen in Fig. 1.




Figure 1: The Franke test function plotted for $0 \leq x, y \leq 1$.

Terrain data

The terrain data used is taken from the U.S. Department of the Interior U.S. Geological Survey’s (USGS) EarthExplorer³ website. The USGS stores data from the Shuttle Radar Topography Mission (SRTM) which maps the earth’s land surface topology with a resolution of 1 arc-second (about 30 m). We will use SRTM data taken from the EarthExplorer website as the basis for our terrain parametrization.

The specific terrain data we will use in the present project is taken from the Møsvatn Austfjell area in the municipality of Tinn in Telemark county, Norway. A visual representation of the data is shown in Fig. 2.

³EarthExplorer website: <https://earthexplorer.usgs.gov/>.

⁴All tests can be run automatically by calling `pytest -v` from anywhere within the Github repository

Table 1: Parameters β and their bootstrap computed variance $\sigma^2(\beta)$ for the OLS fits of the Franke function, shown in Fig. 3. Pair wise columns represent β and $\sigma^2(\beta)$ for each polynomial degree p used. Each row shows the β_j coefficient and $\sigma^2(\beta)$ for the corresponding monomial _{j} term.

p	2		3		4		5	
	β	$\sigma^2(\beta)$	β	$\sigma^2(\beta)$	β	$\sigma^2(\beta)$	β	$\sigma^2(\beta)$
1	1.176 06	0.000 07	0.994 32	0.000 14	0.607 14	0.000 06	0.379 20	0.000 52
x	-1.065 33	0.000 47	-0.603 90	0.003 14	4.231 95	0.004 78	8.057 87	0.039 89
y	-0.741 96	0.000 41	1.367 10	0.002 06	3.188 95	0.003 99	3.773 85	0.033 82
x^2	0.120 32	0.000 26	-1.401 95	0.008 85	-19.382 16	0.055 41	-35.011 95	0.520 45
xy	0.875 01	0.000 25	2.043 06	0.005 97	-2.283 49	0.028 25	-15.642 35	0.500 91
y^2	-0.367 20	0.000 29	-6.659 28	0.005 85	-12.536 72	0.044 34	-8.308 69	0.440 47
x^3			0.891 77	0.002 88	25.256 28	0.107 94	49.218 54	1.617 46
x^2y			0.362 61	0.002 19	8.104 44	0.056 90	45.875 08	1.703 97
xy^2			-1.506 62	0.001 88	1.405 54	0.052 42	21.251 27	1.474 19
y^3			4.696 44	0.002 07	12.615 20	0.084 71	-9.030 64	1.506 07
x^4					-10.841 36	0.025 48	-24.135 09	1.257 84
x^3y					-5.151 14	0.019 48	-54.812 83	1.391 24
x^2y^2					0.007 64	0.014 37	-8.235 81	1.141 34
xy^3					-1.905 24	0.020 13	-30.191 86	1.123 19
y^4					-3.511 79	0.020 32	30.201 66	1.282 60
x^5							1.538 82	0.160 41
x^4y							19.535 04	0.185 67
x^3y^2							10.839 52	0.170 47
x^2y^3							-5.299 38	0.159 35
xy^4							16.912 43	0.150 35
y^5							-16.843 69	0.169 20




Figure 3: Ordinary least squares fits, using data from the Franke function with polynomials of degree 2, 3, 4, and 5. The p parameter indicates what order of polynomials are used. The target function of Franke can be seen in Fig. 1.




Figure 4: The absolute difference between the ordinary least squares fits of Fig. 3 and the true data, the Franke function. Polynomials of degree 2, 3, 4, and 5 have been used in the fitting. The p parameter indicates what order of polynomials are used.




Figure 2: The terrain data in use in the present work, taken from the Møsvatn Austfjell area in the municipality of Tinn in Telemark county, Norway. Retrieved using the USGS EarthExplorer website. The height data is scaled to fit in $0 \leq z \leq 1$, and the reference zero point is set to zero.

RESULTS AND DISCUSSION

Verification of the models: The Franke function

Our first tests comprise simple OLS fitting of the Franke function described in section I. We perform the fits using the scheme described in section B. Basic testing is done automatically using the `pytest` unit test framework.⁴ These ensure e.g. that the fits performed by the manual matrix inversion are identical to the ones performed automatically by `sci-kit learn`, that intercepts and polynomial coefficients are fitted correctly when the underlying functional form is itself a polynomial of known order, etc. In addition to the small scale automatic testing, we perform more thorough, extensive testing on the results of the fittings. First off we visualize and compare the performed OLS fits to the underlying Franke function. These are shown in Fig. 3 (the Franke function, for reference, is shown in Fig. 1). We note that, as expected, higher degree polynomials appear to recreate the underlying exponentials better than lower degree polynomials.

The fact that higher order polynomials are better able to reproduce the Franke function is in itself not surprising. Remember that Franke formed his function from a sum of exponentials. The power series expansion of this function is

thus just a sum over all possible monomials in x and y with prefactors proportional to $1/n!$.

The resampling scheme described in section H, we compute the variance of the computed β coefficients. These are shown in Table 1 for all the different polynomial degrees used, ranging from $p = 2$ to $p = 5$. We note that increasing the polynomial degrees apparently tends to increase the variance of the computed β coefficients. This is in line with the statement of the bias-variance tradeoff: Higher model complexity leads to an increased sensitivity to small fluctuations in the training set, and the variance goes up.

The absolute difference between the fitted models and the Franke function is shown in Fig. 4. We note that the magnitude of the error decreases as the polynomial order increases, but the error for higher order polynomial fits exhibit faster oscillating behaviour.

Terrain data parametrization

Nulla malesuada porttitor diam. Donec felis erat, congue non, volutpat at, tincidunt tristique, libero. Vivamus viverra fermentum felis. Donec nonummy pellentesque ante. Phasellus adipiscing semper elit. Proin fermentum massa ac quam. Sed diam turpis, molestie vitae, placerat a, molestie nec, leo. Maecenas lacinia. Nam ipsum ligula, eleifend at, accumsan nec, suscipit a, ipsum. Morbi blandit ligula feugiat magna. Nunc eleifend consequat lorem. Sed lacinia nulla vitae enim. Pellentesque tincidunt purus vel magna. Integer non enim. Praesent euismod nunc eu purus. Donec bibendum quam in tellus. Nullam cursus pulvinar lectus. Donec et mi. Nam vulputate metus eu enim. Vestibulum pellentesque felis eu massa.

CONCLUSION

Nulla malesuada porttitor diam. Donec felis erat, congue non, volutpat at, tincidunt tristique, libero. Vivamus viverra fermentum felis. Donec nonummy pellentesque ante. Phasellus adipiscing semper elit. Proin fermentum massa ac quam. Sed diam turpis, molestie vitae, placerat a, molestie nec, leo. Maecenas lacinia. Nam ipsum ligula, eleifend at, accumsan nec, suscipit a, ipsum. Morbi blandit ligula feugiat magna. Nunc eleifend consequat lorem. Sed lacinia nulla vitae enim. Pellentesque tincidunt purus vel magna. Integer non enim. Praesent euismod

nunc eu purus. Donec bibendum quam in tel- mi. Nam vulputate metus eu enim. Vestibulum
lus. Nullam cursus pulvinar lectus. Donec et pellentesque felis eu massa.